

Technical Whitepaper

Technical Whitepaper — LLAMA_CPP

Company: Anticloud FZ LLE

Model: PAX L5 Narrow L2 General 27B

Project: LLAMA_CPP | Category: CHIP_QUANTIZATION

Upstream: <https://github.com/ggerganov/llama.cpp> (MIT)

Overview

This document covers technical whitepaper for the Anticloud integration of LLAMA_CPP.

CPU-optimized quantized LLM inference

Anticloud Integration

PAX L5 Narrow L2 General 27B is integrated into LLAMA_CPP to provide:

- Local AI inference with zero cloud dependency
- AIOSS tamper-evident audit chain
- AES-256 encryption at rest
- Single-binary deployment

Contact

Lois-Kleinner Alpasan — CEO & CTO, Anticloud FZ LLE

lois@0-1.gg | 0-1.gg